

struct file_operations Cheat Sheet (Linux Kernel)

This structure defines how your driver responds to file-related system calls like `open()`, `read()`, `write()`, etc. Found in `<linux/fs.h>`, commonly used in character drivers.

Meta Information

struct module *owner

- **Purpose:** Prevents the module from being unloaded while it's in use.
 - **Usage:** Always set to `THIS_MODULE`.
-

File Position Control

```
loff_t (*llseek)(struct file *filp, loff_t offset, int origin);
```

- **Purpose:** Seek to a new file position.
 - **Returns:** New offset (≥ 0), or error (< 0).
 - **If NULL:** Kernel may allow strange seeking behavior.
-

Reading

```
ssize_t (*read)(struct file *filp, char __user *buf, size_t count, loff_t *offset);
```

- **Purpose:** Read data from the device.

- **Returns:** Number of bytes read, or error.
- **If NULL:** `read()` system call returns `-EINVAL`.

```
ssize_t (*aio_read)(struct kiocb *, char __user *,  
size_t, loff_t);
```

- **Purpose:** Asynchronous version of `read`.
 - **If NULL:** Falls back to synchronous `read`.
-

Writing

```
ssize_t (*write)(struct file *filp, const char __user  
*buf, size_t count, loff_t *offset);
```

- **Purpose:** Write data to the device.
- **Returns:** Number of bytes written, or error.
- **If NULL:** `write()` returns `-EINVAL`.

```
ssize_t (*aio_write)(struct kiocb *, const char __user *,  
size_t, loff_t);
```

- **Purpose:** Asynchronous write.
-

Directory Ops (Filesystems Only)

```
int (*readdir)(struct file *, void *, filldir_t);
```

- **Used by:** Filesystems, not device drivers.
- **If NULL:** Should be for device drivers.



Polling

```
unsigned int (*poll)(struct file *filp, struct poll_table_struct *wait);
```

- **Purpose:** Supports `poll()`, `select()`, and `epoll()`.
- **Returns:** Bitmask indicating readiness for I/O.
- **If NULL:** Device assumed readable/writable without blocking.



IOCTL

```
int (*ioctl)(struct inode *, struct file *, unsigned int cmd, unsigned long arg);
```

- **Purpose:** Device-specific commands.
- **Returns:** `-ENOTTY` if unsupported.



Memory Mapping

```
int (*mmap)(struct file *, struct vm_area_struct *);
```

- **Purpose:** Map device memory into user space.
- **If NULL:** `mmap()` returns `-ENODEV`.



Open/Close

```
int (*open)(struct inode *, struct file *);
```

- **Purpose:** Called when file is opened.
- **If NULL:** Open always succeeds silently.

```
int (*release)(struct inode *, struct file *);
```

- **Purpose:** Called when file is closed.
 - **If NULL:** Close succeeds silently.
-

Flushing & Syncing

```
int (*flush)(struct file *);
```

- **Purpose:** Flush data before closing a file descriptor.
- **Rarely used.**

```
int (*fsync)(struct file *, struct dentry *, int datasync);
```

- **Purpose:** Called by `fsync()` to flush pending data.
- **If NULL:** Returns `-EINVAL`.

```
int (*aio_fsync)(struct kiocb *, int datasync);
```

- **Purpose:** Async version of `fsync`.
-

Async I/O & Notification

```
int (*fasync)(int fd, struct file *filp, int mode);
```

- **Purpose:** Asynchronous notification (e.g., SIGIO).

- **Advanced use case.**
-

File Locking

```
int (*lock)(struct file *, int cmd, struct file_lock *);
```

- **Purpose:** File locking.
 - **Rarely used in device drivers.**
-

Vector I/O

```
ssize_t (*readv)(struct file *, const struct iovec *,  
unsigned long, loff_t *);
```

```
ssize_t (*writev)(struct file *, const struct iovec *,  
unsigned long, loff_t *);
```

- **Purpose:** Scatter-gather I/O operations.
 - **If NULL:** Falls back to normal `read/write`.
-

Zero-copy Transfer

```
ssize_t (*sendfile)(struct file *, loff_t *, size_t,  
read_actor_t, void *);
```

- **Purpose:** Efficient file transfer (e.g., for web servers).
- **If NULL:** `sendfile()` fails.

```
ssize_t (*sendpage)(struct file *, struct page *, int  
offset, size_t size, loff_t *pos, int more);
```

- **Purpose:** Kernel uses this to send memory pages to a file descriptor.
-

Memory Area Hints

`unsigned long (*get_unmapped_area)(...)`

- **Purpose:** Helps the kernel find aligned memory regions for mmap.
 - **Rarely implemented.**
-

Miscellaneous

`int (*check_flags)(int flags);`

- **Purpose:** Validate flags in `fcntl(F_SETFL)`.

`int (*dir_notify)(struct file *, unsigned long);`

- **Purpose:** Filesystem-specific — notify on directory changes.
-

Notes

- **NULL values** in most fields mean “not supported.”
- **Return values** follow the standard kernel convention: nonnegative = success, negative = error (`-EINVAL`, `-ENODEV`, etc.).
- Most **device drivers** only implement a subset: `open`, `release`, `read`, `write`, `mmap`, `poll`, `ioctl`.